

Practical Machine Learning Project

Yong-Hao Bai

7/10/2020

Load the required packages

```
library(caret); library(rattle); library(rpart); library(rpart.plot); library(randomForest); library(repmis);  
library(lattice); library(ggplot2); library(readr); library(gbm)
```

Load the Data, divide the data

```
set.seed(19)  
trainurl = "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"  
testurl = "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv"  
download.file(trainurl, "pml-training.csv")  
download.file(testurl, "pml-testing.csv")  
training <- read.csv("pml-training.csv", na.strings=c("NA", "#DIV/0!", ""))  
testing <- read.csv("pml-testing.csv", na.strings=c("NA", "#DIV/0!", ""))  
#update datasets to exclude those variables with NA values  
training <- training[, colSums(is.na(training)) == 0]  
testing <- testing[, colSums(is.na(testing)) == 0]
```

Remove irrelevant variables to the prediction

```
newtraining <- training[, -c(1:7)]  
newtesting <- testing[, -c(1:7)]
```

For cross validation purpose, the training data will be split into training training and training testing.


Data Modeling

Test the predictive power by trying different methods

Decision Tree

Classification Tree

- A
- B
- C
- D
- E



Confusion Matrix and Statistics

##

Reference

## Prediction	A	B	C	D	E
## A	1499	243	23	109	33
## B	54	689	124	71	202
## C	47	107	810	136	127
## D	62	79	69	623	79
## E	12	21	0	25	641

##

Overall Statistics

##

Accuracy : 0.7242

95% CI : (0.7126, 0.7356)

No Information Rate : 0.2845

P-Value [Acc > NIR] : < 2.2e-16

##

Kappa : 0.6495

##

McNemar's Test P-Value : < 2.2e-16

##

Statistics by Class:

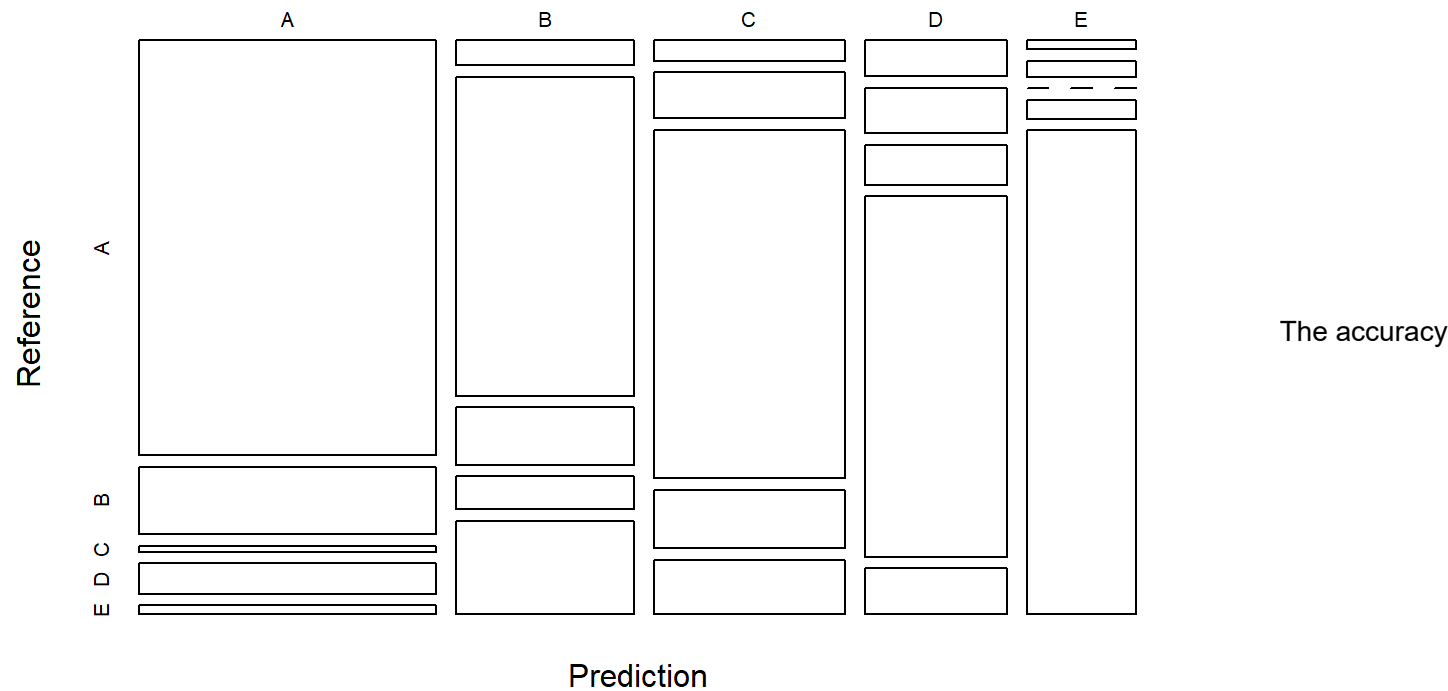
##

##	Class: A	Class: B	Class: C	Class: D	Class: E
## Sensitivity	0.8955	0.6049	0.7895	0.6463	0.5924
## Specificity	0.9031	0.9050	0.9142	0.9413	0.9879
## Pos Pred Value	0.7861	0.6044	0.6601	0.6831	0.9170
## Neg Pred Value	0.9560	0.9052	0.9536	0.9314	0.9150
## Prevalence	0.2845	0.1935	0.1743	0.1638	0.1839
## Detection Rate	0.2547	0.1171	0.1376	0.1059	0.1089
## Detection Prevalence	0.3240	0.1937	0.2085	0.1550	0.1188
## Balanced Accuracy	0.8993	0.7549	0.8518	0.7938	0.7902

Accuracy

0.7242141

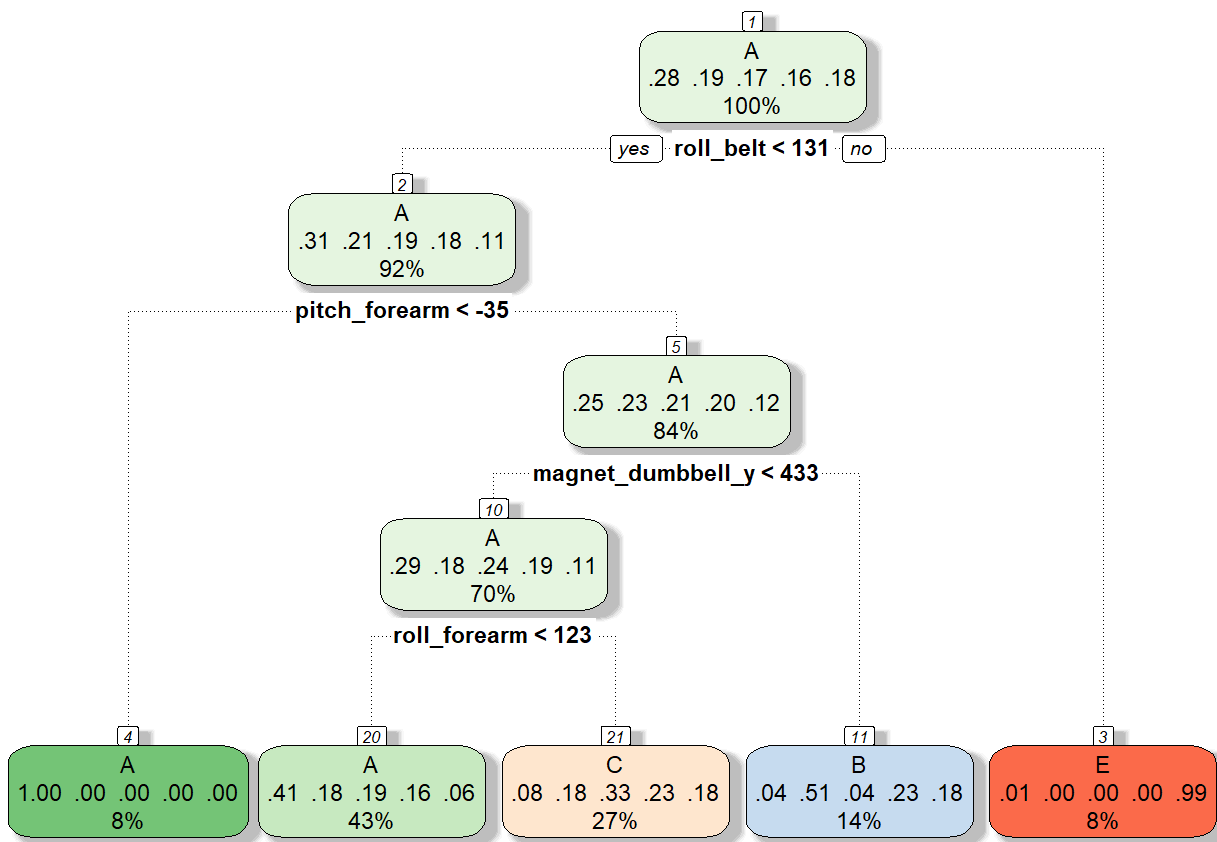
Decision Tree - Accuracy = 0.7242



rate of the model is low: 0.7242.

Classification tree

```
## CART
##
## 13737 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 10990, 10989, 10990, 10989, 10990
## Resampling results across tuning parameters:
##
##  cp      Accuracy  Kappa
##  0.03550  0.5112    0.36227
##  0.06052  0.4414    0.25185
##  0.11688  0.2998    0.02353
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.0355.
```



Rattle 2020-Jul-11 23:33:41 josep

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A     B     C     D     E
##           A 1526   31  112    0    5
##           B  490  379  270    0    0
##           C  476   38  512    0    0
##           D  436  174  354    0    0
##           E  158  157  294    0  473
##
## Overall Statistics
##
##           Accuracy : 0.4911
##           95% CI : (0.4782, 0.5039)
##           No Information Rate : 0.5244
##           P-Value [Acc > NIR] : 1
##
##           Kappa : 0.3344
##
## Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.4945   0.4865   0.3320      NA  0.98954
## Specificity      0.9471   0.8512   0.8816   0.8362  0.88737
## Pos Pred Value   0.9116   0.3327   0.4990      NA  0.43715
## Neg Pred Value    0.6295   0.9157   0.7880      NA  0.99896
## Prevalence       0.5244   0.1324   0.2620   0.0000  0.08122
## Detection Rate    0.2593   0.0644   0.0870   0.0000  0.08037
## Detection Prevalence 0.2845   0.1935   0.1743   0.1638  0.18386
## Balanced Accuracy 0.7208   0.6688   0.6068      NA  0.93845
```

```
## Accuracy
## 0.491079
```

The accuracy rate of the model is even lower.

Boosted Logistic Regression

```
## Boosted Logistic Regression
##
## 13737 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 10989, 10989, 10990, 10991, 10989
## Resampling results across tuning parameters:
##
##   nIter  Accuracy   Kappa
##   11     0.8194597  0.7688882
##   21     0.8704476  0.8344241
##   31     0.8958802  0.8674851
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was nIter = 31.
```

```
## Accuracy
## 0.8967284
```

The accuracy rate of the model has improved from the prior 2 models.

Gradient Boosting

```
## A gradient boosted model with multinomial loss function.
## 150 iterations were performed.
## There were 52 predictors of which 52 had non-zero influence.
```

```
## Stochastic Gradient Boosting
##
## 13737 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 1 times)
## Summary of sample sizes: 10990, 10988, 10991, 10990, 10989
## Resampling results across tuning parameters:
##
##  interaction.depth  n.trees  Accuracy   Kappa
##  1                   50       0.7473965  0.6795041
##  1                   100      0.8189556  0.7707698
##  1                   150      0.8516405  0.8122365
##  2                    50      0.8528048  0.8134939
##  2                   100      0.9060922  0.8811091
##  2                   150      0.9303334  0.9118307
##  3                    50      0.8956102  0.8678282
##  3                   100      0.9430734  0.9279625
##  3                   150      0.9595977  0.9488814
##
## Tuning parameter 'shrinkage' was held constant at a value of 0.1
##
## Tuning parameter 'n.minobsinnode' was held constant at a value of 10
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were n.trees = 150, interaction.depth =
## 3, shrinkage = 0.1 and n.minobsinnode = 10.
```



```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1653   36    0    0    0
##           B   13 1063   30    8   16
##           C    6   38  983   27    8
##           D    1    2   10  923   20
##           E    1    0    3    6 1038
##
## Overall Statistics
##
##           Accuracy : 0.9618
##           95% CI : (0.9565, 0.9665)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9516
##
## Mcnemar's Test P-Value : 1.357e-08
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity       0.9875   0.9333   0.9581   0.9575   0.9593
## Specificity       0.9915   0.9859   0.9837   0.9933   0.9979
## Pos Pred Value    0.9787   0.9407   0.9256   0.9655   0.9905
## Neg Pred Value     0.9950   0.9840   0.9911   0.9917   0.9909
## Prevalence        0.2845   0.1935   0.1743   0.1638   0.1839
## Detection Rate     0.2809   0.1806   0.1670   0.1568   0.1764
## Detection Prevalence 0.2870   0.1920   0.1805   0.1624   0.1781
## Balanced Accuracy  0.9895   0.9596   0.9709   0.9754   0.9786
```

```
## Accuracy
## 0.9617672
```

The Accuracy is getting better.

Random Forest

```
##
## Call:
## randomForest(formula = classe ~ ., data = training_train, method = "class")
##           Type of random forest: classification
##           Number of trees: 500
##           No. of variables tried at each split: 7
##
##           OOB estimate of error rate: 0.48%
## Confusion matrix:
##           A    B    C    D    E class.error
## A 3904     2    0    0    0 0.0005120328
## B   12 2638     8    0    0 0.0075244545
## C    0   15 2381     0    0 0.0062604341
## D    0    0   22 2229     1 0.0102131439
## E    0    0    2    4 2519 0.0023762376
```

Confusion Matrix and Statistics

##

Reference

## Prediction	A	B	C	D	E
## A	1673	6	0	0	0
## B	0	1132	2	0	0
## C	0	1	1022	13	1
## D	0	0	2	950	0
## E	1	0	0	1	1081

##

Overall Statistics

##

Accuracy : 0.9954

95% CI : (0.9933, 0.997)

No Information Rate : 0.2845

P-Value [Acc > NIR] : < 2.2e-16

##

Kappa : 0.9942

##

McNemar's Test P-Value : NA

##

Statistics by Class:

##

##	Class: A	Class: B	Class: C	Class: D	Class: E
## Sensitivity	0.9994	0.9939	0.9961	0.9855	0.9991
## Specificity	0.9986	0.9996	0.9969	0.9996	0.9996
## Pos Pred Value	0.9964	0.9982	0.9855	0.9979	0.9982
## Neg Pred Value	0.9998	0.9985	0.9992	0.9972	0.9998
## Prevalence	0.2845	0.1935	0.1743	0.1638	0.1839
## Detection Rate	0.2843	0.1924	0.1737	0.1614	0.1837
## Detection Prevalence	0.2853	0.1927	0.1762	0.1618	0.1840
## Balanced Accuracy	0.9990	0.9967	0.9965	0.9925	0.9993


Accuracy

0.9954121


Looking at the results, clearly, the random forest model provides a more accurate prediction of classe. The expected out-of-sample error is estimated at 0.005.

plot of the model error rate by number of trees and 20 most important variables (out of 52)

Random forest model error rate by number of trees



Accurary comparison among models



Random Forest has the highest accuracy.

Prediction on Testing


Based on Random Forest prediction:

```
## 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
## B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

Appendix

check for correlation

```
## corrplot 0.84 loaded
```



Variable Importance

##	Overall
## roll_belt	869.54426
## pitch_belt	480.80822
## yaw_belt	634.18014
## total_accel_belt	158.91033
## gyros_belt_x	63.94466
## gyros_belt_y	76.19367
## gyros_belt_z	219.52569
## accel_belt_x	80.27219
## accel_belt_y	82.95947
## accel_belt_z	268.10200
## magnet_belt_x	186.61420
## magnet_belt_y	275.78663
## magnet_belt_z	293.60114
## roll_arm	213.62676
## pitch_arm	124.15713
## yaw_arm	161.48501
## total_accel_arm	67.54977
## gyros_arm_x	96.47692
## gyros_arm_y	96.40537
## gyros_arm_z	44.92191
## accel_arm_x	162.16946
## accel_arm_y	113.43433
## accel_arm_z	94.60188
## magnet_arm_x	183.19255
## magnet_arm_y	163.13118
## magnet_arm_z	129.01364
## roll_dumbbell	293.17185
## pitch_dumbbell	126.78872
## yaw_dumbbell	177.74233
## total_accel_dumbbell	194.09640
## gyros_dumbbell_x	90.94863
## gyros_dumbbell_y	169.50741
## gyros_dumbbell_z	63.29426
## accel_dumbbell_x	177.25603
## accel_dumbbell_y	298.05300
## accel_dumbbell_z	246.44862
## magnet_dumbbell_x	339.16801
## magnet_dumbbell_y	477.30535
## magnet_dumbbell_z	523.25405
## roll_forearm	421.36225
## pitch_forearm	534.13935
## yaw_forearm	116.62986
## total_accel_forearm	75.09184
## gyros_forearm_x	54.18130
## gyros_forearm_y	86.85268
## gyros_forearm_z	56.60624
## accel_forearm_x	219.74014
## accel_forearm_y	101.00889
## accel_forearm_z	166.67266
## magnet_forearm_x	150.19272
## magnet_forearm_y	156.75580
## magnet_forearm_z	203.17271